



UTM
UNIVERSITI TEKNOLOGI MALAYSIA

DEPARTMENT OF COMPUTER SCIENCE

FACULTY OF COMPUTING

UNIVERSITI TEKNOLOGI MALAYSIA

SECP3843 - SPECIAL TOPIC IN DATA ENGINEERING (WBL)

CLOUD ARCHITECTURE IN DATA ENGINEERING:

**A SYSTEMATIC REVIEW OF INFRASTRUCTURE DESIGN, ADOPTION
CHALLENGES, AND FUTURE DIRECTIONS**

PREPARED BY : GROUP 2

NAME	MATRIC NO
AIN NURNABILA BINTI MOHD AZHAR	A23CS0207
NURUL ADRIANA BINTI KAMAL JEFRI	A23CS0258

PREPARED FOR:

DR. ARYATI BINTI BAKRI

Cloud Architecture in Data Engineering: A Systematic Review of Infrastructure Design, Adoption Challenges, and Future Directions

By: Ain Nurnabila, Nurul Adriana

Executive Summary

Cloud architecture now has become the foundational framework of modern data engineering which lets enterprises build scalable pipelines free from the effects of traditional on-premise systems. Architectural patterns like Data Lakehouse have transformed how data is ingested, processed, stored, and served across platforms involving AWS, GCP, and Microsoft Azure. This technical report studies the functional layers of a cloud data engineering stack, explains the advantages and disadvantages of cloud adoption, and identifies evolving trends including multi-cloud strategies and AI-enhanced pipelines. Key recommendations include adopting a Lakehouse architecture, prioritising serverless managed services, and embedding proactive cost monitoring through FinOps practices. As organizations are highly dependent on data-driven decision making, cloud architecture will remain the essential enabler of data engineering.

1. Introduction

The way businesses gather, store, process, as well as analyze data at scale has completely changed as a result of cloud computing. Cloud architecture has become the main framework for creating modern data pipelines in the field of data engineering,

which is in charge of building and managing the architecture that allows transporting data from source to destination. Businesses increasingly use managed cloud services from the providers like Microsoft Azure, AWS, and GCP in order to handle data workloads starting from batch processing to real-time analytics instead of relying on expensive on-premise servers with already fixed capacity.

This study pursues three goals: charting the evolution from traditional data warehouses to modern Lakehouse architectures; examining how cloud platforms hold serverless services for transformation, storage, and serving; and evaluating advantages, challenges, and future directions of cloud adoption. By doing research on literature reviews and real-world case studies, the report proves insights that are relevant to technical and business audiences.

2. Literature Review

The existing cloud architecture has been designed in terms of a data science hierarchy of needs, with basic infrastructure tying together all the downstream analytics. In this framework, the practitioners have advanced to data lifecycle engineers, who control the flow of data in the end-to-end manner (Reis & Housley, 2022). These functions are in contrast to Type A engineers who use high

level abstraction versus Type B engineers who develop custom infrastructure.

Additionally, engineers also serve as organizational middlemen between technical production and business leadership to technical deployment. New paradigms such as the data mesh support this change by encouraging decentralized ownership and data as a product (Machado et al., 2022). To leverage the value of these investments, the C-suite should take the analytics agenda, so that engineering decisions are not isolated to IT departments, but are strategically aligned with business results (Chin et al., 2017).

Undercurrents such as DataOps and financial management are what keep the operations running smoothly and enhance pipeline reliability due to the integration of Agile and Lean approaches (Rodriguez et al., 2020). This infrastructure can be ultimately valuable because it enables data to be made actionable with reverse ETL that transfers enriched insights of warehouses directly back to operational systems (George, 2023). To cope with these intricate settings, metadata-first designs deliver the visibility and provenance that are the fundamentals of sound auditing, control, and organizational adaptability over a long duration (Guntupalli & Venkata, 2021).

3. Analysis and Discussion

3.1 Functional Layers of Cloud Data Engineering




Figure 1 Cloud data engineering stack showing the four primary layers and associated platform tools

As shown in Figure 1, a modern cloud data engineering stack seems to function across four main layers which are intake, storage, processing, and serving. Businesses may build effective data pipelines thanks to the complete data engineering design patterns that major cloud providers like Microsoft Azure, GCP, and AWS have created especially for their platforms (Rucco et al., 2025). Data engineering teams may focus more on pipeline logic and data quality as every layer is supported by specially managed services that abstract architecture complexity.

Tools like AWS Kinesis, Azure Event Hubs, and Google Cloud Pub/Sub gather data at the uptake layer in both batch and real-time modes from a variety of sources, including relational databases,

streaming feeds, Internet of Things (IoT) devices, and Application Programming Interfaces (APIs). In any case of source volume, these solutions serve as the important entry point of the data engineering pipeline, guaranteeing accurate and fast data collecting. Moreover, the delay and performance characteristics of the entire downstream pipeline are often directly chosen by the choice of uptake architecture, whether it be batch-based, stream-based, or hybrid (Khine & Wang, 2024).

Cloud systems provide server-free and managed compute alternatives at the processing layer the scale continuously due to workload demands. With each function carrying out a specific processing task and passing processed data to the next stage, server-free technologies such as AWS Lambda, Google Cloud Functions, and Azure Functions can also be applied across ingestion, processing, and storage stages. This enables flexibility and scalability as functions can be added or removed based on workload requirements (Shojaee Rad & Ghobaei-Arani, 2024). By eliminating the need for manual resource allocation and architecture management, this event-driven framework minimizes operational expenses and implementation time for data engineering teams.

The Data Lakehouse concept, which combines scalability of data lakes with organized search performance of modern data warehouses, is increasingly serving as the foundation of the storage layer. Based on open direct-access data formats for instance Apache Parquet, the Lakehouse architecture

addresses the main issues with traditional data warehouse architectures (Zaharia et al., 2021). Building on this foundation, recent research has proven that a Data Lakehouse platform created with cloud-based ecosystems like Kubernetes and DataOps principles can be made adaptable, able to operate on any cloud provider or on-premise environment while effectively managing millions of records without vendor lock-in (AbouZaid et al., 2025). Platforms which are Google BigQuery, Snowflake, and Azure Synapse Analytics complete the data pipeline lifecycle at the serving layer by delivering processed analytical results to dashboards, BI Tools, and downstream machine learning applications.

3.2 Real-World Case Examples and Comparative Analysis

Dimension	Advantages	Disadvantages
Scalability	Elastic, auto-scales with workload demand	Over-scaling may cause unexpected costs
Cost model	Pay-as-you-go reduces capital expenditure	Data egress fees can accumulate
Operations	Managed services reduce infrastructure overhead	Dependency on vendor uptime and reliability
Flexibility	Supports multi-cloud and hybrid strategies	Proprietary services risk vendor lock-in
Security	Built-in IAM, encryption, and compliance tools	Data sovereignty and GDPR compliance needed

Table 1 *Comparative summary of advantages and disadvantages of cloud architecture in data engineering.*

Widely recognized industry deployments prove how cloud architecture might be used practically in data engineering. The main advantages and disadvantages of adopting cloud architecture, as noted in the examined literature, are summarized in Table 1. With the introduction of 5G networks, IoT technologies, and edge computing creating

increasingly complex data environments which traditional on-premise systems may not be able to transform, process, and visualize the wide range of insights required by modern organizations, big data is steadily growing beyond its foundational constructs of variety as the table indicates (Mills et al., 2024). At the corporate level, cloud architecture immediately reacts to these needs.

One of the most well-known real-world applications of cloud-based data engineering is Netflix. Cloud-based big data analytics offers quantifiable benefits such as enhanced evaluation, improved safety, and streamlined infrastructure according to Abueid (2024). AWS has played an important role in supporting big scale data analytics for Netflix, which heavily depends on analytics to customize content recommendations, optimize streaming quality, and enhance overall user experience.

3.3 Adoption Challenges and Future Trends



Figure 2 Primary adoption challenges in cloud data engineering and corresponding mitigation strategies (Radixweb, n.d.)

The main barriers to cloud adoption include cultural barriers, and data silos that continue to exist, limiting the collaboration flow

needed in modern DataOps (Rodriguez et al., 2020; Patel, 2019). A successful move to these architectures requires a fundamental organizational change of isolated team arrangements to a more essentially cultural data management approach (Reis & Housley, 2022).

This barrier is also exacerbated by a large digital skills gap, since in many traditional educational systems, there is no practical training in distributed systems and real-time streaming needed to make modern-day engineering happen (Ponnekanti, 2025). Table 2 illustrates that the key to such technical and structural issues is to make the data engineers the organizational connectors between fragmented sources to make them a single, actionable truth (Patel, 2019).

In the future, strategic trends include Type A which is the engineering and metadata-first architecture to become more abstract, transparent, and long-term organizational agile (Reis & Housley, 2022; Guntupalli & Venkata, 2021). FinOps practices that reduce cloud waste by rigorously tagging and automatically monitoring it, support this evolution.

The implementations of Data-as-a-product are not only typified by future strategic directions where cloud architectures provide a decentralized ownership and the option to use managed services to reduce heavy lifting or to develop a custom data infrastructure to meet specific competitive requirements (Reis & Housley, 2022). Finally, the success in the

future will be based on the C-suite leadership that will drive the analytics agenda to make sure that technical pipelines are strategically aligned with the overall business goals (Chin et al., 2017).

4. Recommendation

A number of useful suggestions are made for businesses looking to deploy or update their cloud data infrastructure based on the examination of cloud architecture trends, new developments, and adoption issues in data engineering.

4.1 Adopt a Data Lakehouse Architecture

Organizations should have combined separate data lakes and warehouses into one unified Lakehouse platform by opening table formats for example Delta Lake or Apache Iceberg. This will eliminate redundant data copies and inconsistent governance to support batch and streaming workloads from a single storage layer. A model proven to resolve the reliability and data staleness limitations of traditional warehouse systems (Zaharia et al., 2021).

4.2 Prioritise serverless and Managed Services

Where feasible, self-managed infrastructure should be replaced with managed and serverless alternatives. Services like AWS Lambda reduce manual boarding overhead and lower operational costs and implementation time for data engineering teams (Shojaee Rad & Ghobaei-Arani, 2024).

4.3 Embed Proactive Cost Monitoring

Cloud costs can escalate unpredictably once pipelines are not designed with cost awareness. Organizations should integrate FinOps practices from the outset by using tools like AWS Cost Explorer together with tiered storage policies and automated resource scaling to keep cost effectiveness (Plazotta & Klettke, 2024).

5. Conclusion

Cloud architecture has shifted to be the backbone of present-day data engineering, based on the managed service and Data Lakehouse model to separate technical complexity and improve scalability. This is a modular infrastructure that allows organisations to coalesce the divergent streams of data and offer a high-performance foundation to real-time and batch analytical loads. Using serverless and cloud-native trends, companies can handle large amounts of data faster than at any previous time and still have a production-grade infrastructure.

Finally, the worth of such infrastructure lies in breaking the cultural resistance, data silos, and a high digital skills gap with active C-suite leadership and strategic alignment. To be successful, data engineers must be organizational connectors working between technical operations and business goals. With FinOps to achieve cost accountability and metadata-first designs to provide insight into pipelines, organizations can be resilient and strategically aligned to long term innovation objectives.

Reference

Abueid, A. I. (2024). Big data and cloud computing opportunities and application areas.

Engineering, Technology & Applied Science Research, 14(3), 14509–14516.

<https://doi.org/10.48084/etasr.7339>

AbouZaid, A., Barclay, P. J., Chrysoulas, C., & Pitropakis, N. (2025). Building a modern data platform based on the data lakehouse architecture and cloud-native ecosystem. *Discover Applied Sciences*, 7, Article 294. <https://doi.org/10.1007/s42452-025-06545-w>

Chin, J. K., Hagstroem, M., Libarikian, A., & Rifai, K. (2017). *Advanced analytics: Nine insights from the C-suite*. McKinsey & Company.

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/advanced-analytics-nine-insights-from-the-c-suite>

Deshpande, M. (2023). Rise of DataOps: Streamlining data pipelines and workflows for agile data management. *Journal of Artificial Intelligence Machine Learning and Data Science*, 1(4), 338–342. <https://doi.org/10.51219/jaimld/mahesh-deshpande/94>

Galego, N. M. C., Abreu, P., & Novais, P. (2024). Cloud computing for big data analytics.

Procedia Computer Science, 237, 297–304. <https://doi.org/10.1016/j.procs.2024.05.112>

George, J. (2023). Harnessing the power of real-time analytics and reverse ETL: Strategies for unlocking data-driven insights and enhancing decision-making. *World Journal of Advanced Engineering Technology and Sciences*, 10(1), 190–201.

<https://doi.org/10.30574/wjaets.2023.10.1.0247>

- Guntupalli, B., & Venkata, C. (2021). The role of metadata in modern ETL architecture. *International Journal of Artificial Intelligence Data Science and Machine Learning*, 2. <https://doi.org/10.63282/3050-9262.ijaidsm-l-v2i3p106>
- Khine, P. P., & Wang, Z. (2024). Data lakes: A survey of concepts and architectures. *Computers*, 13(7), Article 183. <https://doi.org/10.3390/computers13070183>
- Machado, I. A., Costa, C., & Santos, M. Y. (2022). Data mesh: Concepts and principles of a paradigm shift in data architectures. *Procedia Computer Science*, 196, 263–271. <https://doi.org/10.1016/j.procs.2021.12.013>
- Mills, N., Issadeen, Z., Matharaarachchi, A., et al. (2024). A cloud-based architecture for explainable big data analytics using self-structuring artificial intelligence. *Discover Artificial Intelligence*, 4, Article 33. <https://doi.org/10.1007/s44163-024-00123-6>
- Patel, J. (2019). Bridging data silos using big data integration. *International Journal of Database Management Systems*, 11(3), 01–06. <https://doi.org/10.5121/ijdms.2019.11301>
- Plazotta, M., & Klettke, M. (2024). Data architectures in cloud environments. *Datenbank-Spektrum*, 24, 243–247. <https://doi.org/10.1007/s13222-024-00490-5>
- Radixweb. (n.d.). *Cloud computing challenges* [Infographic]. <https://radixweb.com>
- Reis, J., & Housley, M. (2022). *Fundamentals of data engineering*. O'Reilly Media.
- Rodriguez, M., De Araújo, L. J. P., & Mazzara, M. (2020). Good practices for the adoption of DataOps in the software industry. *Journal of Physics: Conference Series*, 1694(1), Article 012032. <https://doi.org/10.1088/1742-6596/1694/1/012032>

Rucco, C., Longo, A., & Saad, M. (2025). Efficient data ingestion in cloud-based architecture: A data engineering design pattern proposal. *Data Science and Engineering*.

<https://doi.org/10.1007/s41019-025-00300-2>

Shojaee Rad, Z., & Ghobaei-Arani, M. (2024). Data pipeline approaches in serverless computing: A taxonomy, review, and research trends. *Journal of Big Data*, 11, Article 82.

<https://doi.org/10.1186/s40537-024-00939-0>

Zaharia, M., Ghodsi, A., Xin, R., & Armbrust, M. (2021). Lakehouse: A new generation of open platforms that unify data warehousing and advanced analytics. *Proceedings of the 11th Conference on Innovative Data Systems Research (CIDR)*.

https://www.cidrdb.org/cidr2021/papers/cidr2021_paper17.pdf

Logbook

